

✔ Detailed Answers & Notes

1 DSA — Sliding Window (2 Problems)

Problem A (Fixed Window): Maximum Average Subarray I

LeetCode 643: <https://leetcode.com/problems/maximum-average-subarray-i/>

Hints

- Maintain a sum of the first k elements.
- Slide by: $\text{sum} += a[i] - a[i-k]$.
- Track max sum (or max average = maxSum / k).

C++ Solution

```
#include <bits/stdc++.h>
```

```
using namespace std;
```

```
double findMaxAverage(vector<int>& nums, int k) {
```

```
    long long window = 0;
```

```
    for (int i = 0; i < k; ++i) window += nums[i];
```

```
    long long best = window;
```

```
    for (int i = k; i < (int)nums.size(); ++i) {
```

```
        window += nums[i] - nums[i - k];
```

```
        best = max(best, window);
```

```
    }
```

```
    return (double)best / k;
```

```
}
```

```
int main(){
```

```
    vector<int> v{1,12,-5,-6,50,3};
```

```
    int k = 4;
```

```
    cout << fixed << setprecision(5) << findMaxAverage(v, k); // 12.75000
```

```
}
```

Complexity: $O(n)$ time, $O(1)$ space

Edge Cases: $k == n$, negatives in array, $n == k == 1$

Common Mistakes: Using double for rolling sum (overflow/precision); recomputing sums from scratch.

Problem B (Variable Window): Fruit Into Baskets

LeetCode 904: <https://leetcode.com/problems/fruit-into-baskets/>

Hints

- Maintain a window with **at most 2 distinct** fruit types.
- Use hashmap <type -> count>.
- Shrink from left when size of map > 2.

C++ Solution

```
#include <bits/stdc++.h>
```

```
using namespace std;
```

```
int totalFruit(vector<int>& fruits) {  
    unordered_map<int,int> cnt;  
    int left = 0, best = 0;  
  
    for (int right = 0; right < (int)fruits.size(); ++right) {  
        cnt[fruits[right]]++;  
        while ((int)cnt.size() > 2) {  
            if (--cnt[fruits[left]] == 0) cnt.erase(fruits[left]);  
            ++left;  
        }  
        best = max(best, right - left + 1);  
    }  
    return best;  
}
```

```
int main(){  
    vector<int> f{1,2,1,2,3,2,2};  
    cout << totalFruit(f); // 5 (window [2,1,2,3,2] or [1,2,3,2,2] depending on input)
```

```
}
```

Complexity: $O(n)$ time, $O(1)$ or $O(k)$ space ($k \leq 2$)

Edge Cases: All same type; alternating 2 types; a third type appears once.

Common Mistakes: Not erasing type when count hits 0; shrinking only one step instead of until $\text{map.size()} \leq 2$.

2 SQL — GROUP BY, HAVING, Window Functions

Schema (example)

Employees(emp_id, emp_name, dept_id, salary)

Departments(dept_id, dept_name)

Sales(emp_id, product, region, amount, sale_date)

A) GROUP BY with Multiple Columns

Goal: total sales by region & product.

Stepwise

1. Base aggregation:

```
SELECT region, product, SUM(amount) AS total_amount
```

```
FROM Sales
```

```
GROUP BY region, product;
```

2. Add department dimension (via join):

```
SELECT s.region, s.product, d.dept_name, SUM(s.amount) AS total_amount
```

```
FROM Sales s
```

```
JOIN Employees e ON e.emp_id = s.emp_id
```

```
JOIN Departments d ON d.dept_id = e.dept_id
```

```
GROUP BY s.region, s.product, d.dept_name;
```

Optimization Tips

- Index join keys: Sales(emp_id), Employees(dept_id), Departments(dept_id).
 - If filtering by date/region, add composite index like Sales(region, product, sale_date).
-

B) HAVING (filter groups)

Q: Regions where **sum(amount) > 1,000,000**:

```
SELECT region, SUM(amount) AS total_amount
```

```
FROM Sales
```

GROUP BY region

HAVING SUM(amount) > 1000000;

Use WHERE for row-level filters (e.g., sale_date >= '2025-01-01'), HAVING for aggregated filters.

C) Window Functions (Ranking within partitions)

1. **Top seller per region** (keep ties with RANK):

```
SELECT emp_id, region, SUM(amount) AS total_amount,  
       RANK() OVER (PARTITION BY region ORDER BY SUM(amount) DESC) AS rnk  
FROM Sales  
  
GROUP BY emp_id, region  
  
QUALIFY rnk = 1; -- In Snowflake/BigQuery; in others, wrap as subquery/CTE
```

Portable version (CTE):

```
WITH agg AS (  
    SELECT emp_id, region, SUM(amount) AS total_amount  
    FROM Sales  
    GROUP BY emp_id, region  
)  
  
ranked AS (  
    SELECT emp_id, region, total_amount,  
           RANK() OVER (PARTITION BY region ORDER BY total_amount DESC) AS rnk  
    FROM agg  
)  
  
SELECT * FROM ranked WHERE rnk = 1;
```

2. **Unique row per dept (latest salary record)** with ROW_NUMBER:

```
WITH enriched AS (  
    SELECT e.emp_id, e.emp_name, e.dept_id, e.salary, e.updated_at,  
           ROW_NUMBER() OVER (PARTITION BY e.emp_id ORDER BY e.updated_at DESC) AS rn  
    FROM Employees e  
)  
  
SELECT * FROM enriched WHERE rn = 1;
```

3. **Dense ranking by salary inside department:**

```
SELECT emp_id, dept_id, salary,  
       DENSE_RANK() OVER (PARTITION BY dept_id ORDER BY salary DESC) AS dense_rank_in_dept  
FROM Employees;
```

Optimization Tips

- Partition keys should be selective; ensure useful indexes on join + filter columns.
 - Avoid SELECT * in large windows; project only needed columns.
-

3 CS Core — Notes + 5 Q&A Each (Fully Detailed)

A) Operating Systems — Memory (Paging/Segmentation/VM/TLB/Thrashing)

Short Notes

- **Paging**: fixed-size frames/pages → eliminates external fragmentation.
- **Segmentation**: variable-size logical units (code/data/stack) → external fragmentation possible.
- **Virtual Memory**: logical > physical; demand paging loads on use.
- **TLB**: small associative cache for page table entries.
- **Thrashing**: excessive page faults → CPU idle, disk busy.

Q&A

1. Paging vs Segmentation?

- Paging splits **physical memory** into equal frames; process split into equal pages → easy allocation, no external fragmentation; internal fragmentation possible in last page.
- Segmentation splits program into **logical segments** (varying sizes) → supports protection/sharing; can cause external fragmentation; needs compaction.

2. What is a Page Fault? Steps?

- Access to a page not in RAM. Steps: trap to OS → find a free frame (or evict via replacement) → read page from disk → update page table → restart instruction.

3. Role of TLB?

- Caches recent virtual→physical translations; **hits** avoid page-table walk → huge speedup. TLB miss triggers page-table access.

4. Thrashing — why & prevention?

- Cause: working set > available frames → constant evictions.
- Fix: increase RAM/frames; working-set model; load control; better replacement (e.g., LRU approximation).

5. Replacement policies (FIFO/LRU/Optimal)?

- **FIFO**: simple, Belady's anomaly possible.
 - **LRU**: approximates recency; needs hardware/software support.
 - **Optimal**: clairvoyant benchmark; not implementable.
-

B) DBMS — Normalization & Denormalization

Short Notes

- **1NF**: atomic columns, no repeating groups.
- **2NF**: 1NF + no partial dependency on a composite key.
- **3NF**: 2NF + no transitive dependency on key.
- **BCNF**: for all FDs $X \rightarrow Y$, X must be a superkey (stronger than 3NF).
- **Denormalization**: deliberate redundancy for performance (e.g., analytics).

Q&A

1. Why normalize?

- Reduce redundancy, update anomalies; improves integrity.

2. 3NF vs BCNF?

- 3NF allows $X \rightarrow Y$ when Y is prime attribute; BCNF requires X be a superkey always \rightarrow BCNF stricter.

3. When denormalize?

- Read-heavy analytics, frequent joins causing latency; use materialized views, summary tables.

4. Functional dependency & candidate key?

- FD $X \rightarrow Y$: X determines Y .
- Candidate key: minimal attribute set that functionally determines all attributes.

5. Example of transitive dependency

- Student(roll, dept, dept_hod) with $\text{dept} \rightarrow \text{dept_hod}$. Remove dept_hod to a Department table.
-

C) Computer Networks — DNS, DHCP, ARP

Short Notes

- **DNS**: name \rightarrow IP resolution (root \rightarrow TLD \rightarrow authoritative).
- **DHCP**: dynamic IP allocation (DORA).

- **ARP:** IP→MAC mapping within LAN.

Q&A

1. DNS: iterative vs recursive?

- **Recursive:** resolver asks upstream servers to fetch final answer.
- **Iterative:** resolver queries each server in turn using referrals.

2. DHCP DORA?

- **Discover → Offer → Request → Acknowledge** (broadcast→unicast sequences).

3. What is ARP cache?

- Local table mapping IP to MAC; entries time out; misses cause ARP broadcast.

4. ARP spoofing & mitigation?

- Attacker poisons cache with fake MAC; mitigate using **Dynamic ARP Inspection**, static ARP, switch security.

5. DNS caching benefits & risks

- Faster lookups; but stale records, cache poisoning risk → DNSSEC.

D) OOPs — Encapsulation, Ctors/Dtors (C++)

Short Notes

- **Encapsulation:** bundle data+methods; use access specifiers.
- **Constructors:** default, parameterized, copy, move; **Destructor:** cleanup.
- **RAII:** resource acquisition is initialization.

Q&A

1. Encapsulation benefits?

- Invariants, controlled access, easier refactor; example: getters/setters validate input.

2. Copy vs Move ctor?

- **Copy:** deep copy expensive; **Move:** steals resources (rvalues) → zero-cost transfers.

3. When is destructor called?

- Object lifetime ends (scope exit, delete); free resources (files, sockets, heap).

4. Rule of 3/5/0?

- If you define any of dtor/copy-ctor/copy-assign → define all three (Rule of 3). With move operations → Rule of 5. Prefer compositions with no manual mgmt → Rule of 0.

5. Friend classes — when?

- Tight coupling for performance or operator overloading; use sparingly.
-

E) SDLC — Waterfall vs Agile, Scrum

Short Notes

- **Waterfall:** sequential, rigid requirements.
- **Agile:** iterative, feedback-driven; Scrum ceremonies.

Q&A

1. **When prefer Waterfall?**
 - Stable requirements, regulated domains, heavy documentation.
 2. **Scrum roles & ceremonies?**
 - Roles: Product Owner, Scrum Master, Dev Team.
 - Ceremonies: Planning, Daily Standup, Review, Retrospective.
 3. **Definition of Done vs Acceptance Criteria**
 - DoD: quality checklist across stories; AC: story-specific conditions.
 4. **Agile pros/cons**
 - Pros: adaptability, visibility. Cons: scope creep risk, discipline required.
 5. **Velocity & capacity**
 - Velocity: past completed points; capacity: forecasted points this sprint.
-

F) Software Testing — Unit, Smoke, Regression, Test Pyramid

Short Notes

- **Unit:** test smallest units, fast & isolated.
- **Smoke:** quick sanity on core paths after build.
- **Regression:** ensure changes don't break old features.

Q&A

1. **Smoke vs Sanity?**
 - Smoke: broad, shallow post-build. Sanity: narrow, focused post-fix.
2. **Mocks vs Stubs?**
 - Stubs return canned responses; mocks verify interactions/behavior.
3. **Test Pyramid?**
 - More unit tests, fewer integration, even fewer UI tests → fast + stable.

4. Regression suite selection

- Impact analysis, risk-based prioritization, critical user journeys.

5. Code coverage caveat

- High coverage \neq good tests; assert meaningful behavior.
-

4 Aptitude — Stepwise

Q1 (Time & Work): A in 12 days, B in 18 days. Together?

- Rates: $A=1/12$, $B=1/18 \rightarrow \text{LCM } 36 \rightarrow A=3u$, $B=2u \rightarrow \text{together } 5u = 5/36 \text{ work/day}$.
- Days = $36/5 = 7.2$ days.

Q2 (Pipes): Inlet fills in 6 h, outlet empties in 9 h. Together?

- Net rate = $+1/6 - 1/9 = (3-2)/18 = 1/18 \text{ tank/hour} \rightarrow \mathbf{18 \text{ h}}$.

Q3 (Ratio): A:B = 2:3, B:C = 4:5. Find A:B:C

- Make B equal: $\text{LCM}(3,4)=12 \rightarrow A=8$, $B=12$, $C=15 \rightarrow \mathbf{8:12:15}$.

Shortcuts

- Time & Work: Use LCM to get integer “work units”.
 - Pipes: Convert to net rate (in/out signs).
 - Ratios: Equalize common term via LCM.
-

5 Behavioral — STAR

Q1: Worked under pressure

- **S:** Sprint end; prod bug degrading dashboard latency.
- **T:** Fix within 6 hours without downtime.
- **A:** Rolled back heavy query, added read-replica + cache on hot endpoints; wrote throttling rule; paired with QA for regression.
- **R:** Latency \downarrow **48%**, incident resolved in **3.5h**, post-mortem actions adopted team-wide.

Q2: Motivation in tech

- **S/T:** I enjoy shipping tools that make decisions faster.
- **A:** I keep a rhythm of learning (DSA daily, CS core rotation), contribute to PrepToShine logs, and turn ideas into PoCs.
- **R:** Built SAR colorization demo & faculty dashboard; measurable improvements (accuracy +13%, review time -40%).
Hook: I like the loop: learn \rightarrow build \rightarrow measure \rightarrow iterate.

6 Projects/Resume — Actionables (do today)

- **SAR Colorization:** Add metric bullet —
“Reduced manual review time by **40%** via auto-colorization & QC flags (SSIM + histogram checks).”
- **Faculty Dashboard:** Add metric bullet —
“Cut reporting latency by **55%** using pre-aggregations + caching (Redis).”
- **Resume header:** Add GitHub link to **PrepToShine** and keep daily commit streak.